

EMC Infrastructure for Microsoft Private Cloud

EMC VNX5300, Replication Manager, Microsoft Hyper-V, Microsoft Exchange, SharePoint, SQL Server, System Center

- Optimize infrastructure performance
- Cloud-ready infrastructure
- Automate and simplify management and monitoring

EMC Solutions Group

Abstract

This white paper presents a solution that explores the scalability and performance for mixed application workloads on a Microsoft Hyper-V virtualized platform using an EMC® VNX5300 storage array. It also highlights the ease of management with Microsoft System Center Operations Manager integrated with EMC System Center Management Packs and EMC Storage Integrator.

March 2012



Copyright © 2011-2012 EMC Corporation. All Rights Reserved.

EMC believes the information in this publication is accurate as of its publication date. The information is subject to change without notice.

The information in this publication is provided “as is.” EMC Corporation makes no representations or warranties of any kind with respect to the information in this publication, and specifically disclaims implied warranties of merchantability or fitness for a particular purpose.

Use, copying, and distribution of any EMC software described in this publication requires an applicable software license.

For the most up-to-date listing of EMC product names, see EMC Corporation Trademarks on EMC.com.

Brocade is trademark of Brocade Communications Systems, Inc., in the United States and/or in other countries.

Intel and Xeon are trademarks of Intel Corporation in the U.S. and/or other countries.

All trademarks used herein are the property of their respective owners.

Part Number H8906.1

Table of contents

Executive summary	7
Business case	7
Solution overview	7
Key results	7
Introduction	9
Purpose	9
Scope	9
Audience.....	9
Terminology	9
Overview of components	11
Overview	11
EMC VNX family of unified storage platforms.....	11
EMC VNX5300.....	12
EMC System Center Management Packs.....	12
EMC Storage Integrator	12
EMC Replication Manager	12
Brocade FCX Series IP switches	13
Microsoft Windows Server 2008 R2 with Hyper-V	13
Microsoft System Center Virtual Machine Manager.....	13
Microsoft System Center Operations Manager.....	13
Solution architecture and design	14
Solution architecture.....	14
Environment profile.....	14
Hardware environment.....	15
Software environment	16
Virtual machine configuration	16
Component design and configuration	17
Overview.....	17
Hyper-V cluster design and configuration	17
Virtual machine deployment design	17
Hyper-V network design.....	20
IP switch configuration.....	20
Exchange Server 2010	22
Exchange Server 2010 storage design for the EMC VNX5300	22
Exchange Server 2010 storage design methodology – Building block.....	23
Hyper-V virtual machine design for Exchange Server 2010.....	24

Storage design verification – Jetstress performance test.....	25
SQL Server 2008 R2	26
Storage design overview.....	26
Environment characteristics	26
Design the storage architecture based on environment characteristics.....	26
SQL Server database design overview	27
SharePoint Server 2010	28
Mid-sized SharePoint 2010 design overview	28
Storage design and SharePoint considerations	28
SharePoint 2010 farm Hyper-V and virtual machine design	29
SharePoint farm configuration	30
SharePoint search configuration.....	30
Design consideration for tempdb	31
SCOM.....	31
EMC System Center Management Packs	33
Microsoft SCVVM	35
ESI	37
Data replication and recovery design	40
Overview.....	40
Hyper-V considerations	40
SharePoint farm considerations	40
Exchange considerations	41
SQL Server considerations	42
Test methodology	44
Overview.....	44
Microsoft Exchange Load Generator	44
SQL Server TPC-E like workload	44
SharePoint 2010 VSTS-generated custom workload	44
VSTS test client and test mechanism	45
SharePoint user profiles	45
Performance test results.....	46
Overview.....	46
Hyper-V root server	46
IP switch	46
Exchange Server 2010	47
Exchange performance results.....	47
Storage performance results of Exchange database disks	48
SQL Server 2008 R2	49
Storage performance results.....	49

SQL Server results	50
SharePoint Server 2010	51
SharePoint Performance Results.....	51
Storage performance results of SharePoint.....	52
Failure simulation test results.....	54
Test objective.....	54
Switch failure simulation test.....	54
Overview	54
Exchange Server 2010 impact	54
SQL Server 2008 R2 impact	55
SharePoint Server 2010 impact	56
Hyper-V node failure test.....	57
Overview	57
Exchange Server 2010 impact	59
SharePoint Server 2010 impact	60
VNX5300 SP failure test	61
Overview	61
VNX5300 impact	62
Exchange Server 2010 impact	63
SQL Server 2008 R2 impact	64
SharePoint Server 2010 impact	65
Disk failure test for SQL Server	66
Overview	66
SQL Server 2008 R2 impact	66
Replication Manager test results.....	69
Overview.....	69
Exchange Server 2010	69
Exchange database restore	70
SQL Server 2008 R2	70
Performance impact on the running workload.....	70
Snapshot restore of SQL Server database.....	71
SharePoint Server 2010	72
SharePoint Full Farm replication	72
Content database restore	74
Conclusion	75
Summary	75
Findings.....	75

References.....	77
White papers	77
Product documentation.....	77
Other documentation	77
Supporting information	78
Disable SCSI filtering	78

Executive summary

Business case

Today, many organizations have made a decisive move to revamp the existing storage strategy by running critical applications like Microsoft Exchange Server, SharePoint Server, and SQL Server on a virtualized infrastructure. By consolidating application servers on a virtualized platform, customers can achieve significant cost reductions and increase the environment's ability to scale.

At the same time, it is a constant critical business challenge for IT departments to maintain or improve the performance of a company's mixed Microsoft applications, while providing an easy-to-manage environment. This solution provides a simplified architecture to host different business applications, ensuring that each business line's information is separated from that of the others. It greatly simplifies the environment and reduces operational and management costs. In accordance with the best practices for Microsoft applications, this solution also showcases a comprehensive design methodology to run consolidated workloads across the EMC® VNX5300 storage platform powered by the Intel® Xeon® processor.

Solution overview

There is a growing need among customers to run multiple workloads/applications on a shared infrastructure and meet expected performance levels at lower costs as dictated by the business service-level agreement (SLA). This solution shows a mixed Microsoft workload of Exchange Server 2010, SharePoint Server 2010, SQL Server 2008 R2, and features an element of high availability (HA) in all application environments.

Furthermore, this solution architecture includes the following components to demonstrate a private cloud solution for customers who are looking for enterprise consolidation with management simplicity:

- Microsoft System Center Virtual Machine Manager (SCVMM) and System Center Operations Manager (SCOM) with EMC System Center Management Packs to manage and monitor the whole environment
- Different Microsoft application workloads running on the Hyper-V platform with VNX5300 integrated with EMC Storage Integrator (ESI) for easy storage provisioning to the platform
- Protection of application data provided by EMC Replication Manager using SnapView™ snapshots
- Brocade® FCX Series network switches delivering high performance and low latency network connectivity for both the iSCSI based IP SAN and end user access to the applications services running in virtual machines.

Key results

The solution offers the following key benefits:

- Easy-to-use and simple management features for administrators to provision and manage the infrastructure. It saves 13 steps to create CSV in a windows cluster by integrating with EMC Storage Integrator.
- Sizing guidance for Microsoft Exchange Server, SharePoint Server, and SQL Server virtual environments for mid-sized environments.

- Excellent performance results achieved during the combined workload of all Microsoft applications for:
 - 2,500 concurrent Exchange users, with a 2 GB mailbox size and 0.20 IOPS user profile.
 - 16,440 SharePoint users with 10% user concurrency on the virtualized SharePoint farm.
 - 45,000 users configured for a SQL TPC-E environment with sustained high disk utilization, considering acceptable user response time and saved storage capacity.
- Minimal performance impact during catastrophic, component-level, hardware failure.
- Robust high performance and low latency IP networking from Brocade, that is easy to configure, manage, and monitor.
- Protection of all three Microsoft applications through Replication Manager with SnapView snapshots, with minimal impact on the production environment. The VNX SnapView snapshots job was completed successfully to protect a 1 TB SharePoint farm:
 - Over 3.5 TB Exchange databases
 - Over a 500 GB SQL Server database
- A typical SQL and SharePoint content databases and Exchange database restore takes only a few minutes.
 - The restore of five Exchange databases, including log files (1.8 TB in total), took around 13 minutes. A 250 GB SQL Server database restore took less than four minutes.
 - Five SharePoint content databases (1 TB in total) restore took 7 minutes, 28 seconds.

Introduction

Purpose

The purpose of this document is to describe a validated reference architecture and provide design guidelines for a mixed Microsoft application solution, including Exchange, SQL, and SharePoint Servers, on the EMC VNX5300 storage system. Microsoft Hyper-V is used as the hypervisor platform and the hosts are connected to the storage via iSCSI, with a cost-effective connectivity infrastructure.

Scope

The scope of this paper is to describe:

- The design methodology and considerations for Microsoft applications on a Hyper-V and VNX5300 platform
- Performance testing methodology and test results
- The impact of a hardware failure on Microsoft applications
- The use of Replication Manager to manage the backup and the design considerations
- The performance impact on applications when using Replication Manager snapshots and the instant restore of the applications
- Easy storage provisioning using ESI
- How discovery and health monitoring of the storage environment is possible through SCOM and the EMC System Center Management Packs

Audience

The intended audience for the white paper is:

- Customers
- EMC partners
- Internal EMC personnel

Terminology

Table 1 provides a description of terminology used in this paper.

Table 1. Terminology

Term	Definition
Background Database Maintenance (BDM)	The process of Exchange 2010 database maintenance that involves check summing both active and passive database copies.
Building block	A building block represents the amount of disk and server resources required to support a specified number of Exchange 2010 users. The required resources depend on: <ul style="list-style-type: none">• A specific user profile type• Mailbox size• Disk requirements

Term	Definition
CSV	<p>Without CSV, a failover cluster allows a given disk (LUN) to be accessed by only one node at a time. Given this constraint, each Hyper-V virtual machine in the failover cluster requires its own set of logical units (LUNs) in order to be migrated or fail over independently of other virtual machines.</p> <p>In contrast, on a failover cluster that uses CSV, multiple virtual machines that are distributed across multiple cluster nodes can all access their Virtual Hard Disk (VHD) files at the same time, even if the VHD files are on a single disk (LUN) in the storage. The clustered virtual machines can all fail over independently of one another.</p>
Database availability group (DAG)	A DAG is the base component of the HA and site resilience framework built into Microsoft Exchange Server 2010. A DAG is a group of up to 16 Mailbox servers that hosts a set of databases and provides automatic database-level recovery from failures that affect individual servers or databases.
Pass-through disk	A pass-through disk is where virtual machines have direct access to disks. It is only applicable to block devices such as iSCSI or Fibre Channel (FC).
Recovery time objective (RTO)	RTO is the period of time within which systems, applications, or functions must be recovered after an outage. This defines the amount of downtime that a business can endure.
Virtual Hard Disk (VHD)	This is a publicly available image format specification that allows encapsulation of the hard disk into an individual file for use by the operating system as a virtual disk, in all the same ways that physical hard disks are used. These virtual disks are capable of hosting native file systems (NTFS, FAT, exFAT, and UDFS) while supporting standard disk and file operations.
Volume Shadow Copy (VSS)	<p>The Volume Shadow Copy Service in Windows Server 2008 provides an infrastructure that enables third-party storage management programs, business programs, and hardware providers to cooperate to create and manage shadow copies.</p> <p>VSS coordinates communication between VSS requestors (for example, backup applications), VSS writers (for example, the Exchange 2010 VSS Writer), and VSS providers (system, software, or hardware components that create the shadow copies).</p>